

AUDIT YOUR OWN WEEK 5 PROJECT

1. KICK-OFF (SYSTEM BRIEF)

1. What does the system do?

The system automatically processes incoming vehicle component test reports sent by email. It reads the report, identifies issues or defects mentioned in it, analyzes attached photos, and compares the findings against an existing issue database. If the issue already exists, the system updates the existing record with the new findings. If it is a new issue, it creates a new entry and report. At the end of the process, relevant stakeholders are notified through Slack.

2. What inputs does it take?

The system is not intended to process personal data. The inputs mainly consist of technical information related to vehicle components and testing.

- Emails sent to a dedicated mailbox containing:
 - Markdown (.md) test reports describing component test findings
 - Photos showing the detected issues
- Existing issue records stored in Google Drive folders and reports

3. What does it output?

The system produces:

- AI-generated descriptions of each issue based on the report and photos, which contain its possible causes, consequences, and preventive measures
- Updated or newly created issue reports in the database
- Notifications sent to relevant stakeholders through Slack

4. Who is affected by the output? Who would use it in production?

The system would likely be used by:

- Engineering teams (development and testing)
- Production and quality assurance staff
- External service providers (testing or analysis institutions) or suppliers

The system does not directly affect customers or members of the public.

5. Does a human review the output before any action is taken?

No human review is enforced at the development stage at the end of Week 5. However, in practice, human stakeholders would likely review the database updates after receiving the Slack notification informing that a new report has been processed. They remain responsible for interpreting findings, validating recommendations, and deciding on corrective or preventive actions.

6. Who built it?

The system was built as a custom workflow using n8n and third-party services such as OpenAI, Google Drive/Google Cloud, and Slack APIs. It was thought to be developed by the automotive engineers themselves or commissioned to an internal IT-support team rather than purchased as a standalone commercial product.

2. CFU CHECKPOINTS

This system would most likely be classified as a minimal-risk AI system under the EU AI Act. The workflow automates internal engineering processes and therefore does not fall within any prohibited practices under Article 5, which mainly focuses on the misuse of personal data and unacceptable interactions of the system with the public.

Although the system analyzes technical defects in vehicle components, which are EU-regulated products, its application is limited to internal reporting and knowledge management. The system is not incorporated into the vehicles as a direct safety-critical feature or used for the automated conformity assessment of products. Therefore, it is not classified as category A high-risk. It also does not fall into any of the specific areas in category B (Annex III).

Given that, it is recommended to follow some of the obligations for high-risk AI providers when developing and implementing this system, such as record-keeping and logging (Article 12), human supervision and overriding capabilities (Article 14), or accuracy, robustness and cybersecurity (Article 15).

The use of generative AI to create summaries, recommendations, and classifications could trigger limited transparency considerations under Article 50 if users are not aware that content is AI-generated. However, since the outputs are intended for trained internal staff in a professional workflow and not for interaction with the general public, the practical obligations are likely limited to internal disclosure and documentation rather than formal transparency measures toward consumers.

A fitting possibility for an “AI-generated” label would be the “last modification” field in the output reports, which currently contains only the date and could be extended with either “AI-generated” or the name of the employee who changed it last. Similarly, the Slack notifications could explicitly prompt the users to review the AI-generated content before using it for important technical decisions.

In conclusion, the system fits best into the minimal-risk category, while still benefiting from basic governance measures such as human oversight, logging, and validation of AI-generated findings.

The users of this tool at the current complexity level are likely to be its providers (developers) and deployers at the same time. In the case the adoption was high, and it was planned to extend its use to other departments, the support from an internal IT team or the outsourcing of its development and maintenance could be necessary, giving those the role of providers and the consultants or the engineers themselves becoming the deployers.